

Machine Learning: Using regression and neural networks to fit continuous functions and classify data



Håkon Olav Torvik, Vetle Vikenes & Sigurd Sørli Rustad

FYS-STK4155 – Applied Data Analysis and Machine Learning

Autumn 2021

Department of Physics

University of Oslo

November 10, 2021

ABSTRACT: Abstract coming soon.

Contents

1	Introduction	1
2	Theory	1
2.1	Gradient Decent	1
2.1.1	Ordinary Gradient Decent	2
2.1.2	Stochastic Gradient Decent	2
2.1.3	Adding Momentum	3
2.2	Logistic Regression	3
2.3	Feed-Forward Deep Neural Networks	4
2.3.1	Architecture of Neural Networks	4
2.3.2	Activation Functions	6
2.3.3	Cost Function and Regularization	6
2.3.4	The Backpropagation Algorithm	8
2.3.5	Initialization of weights	9
3	Methods	10
3.1	Franke Function	10
3.1.1	Stochastic Gradient Descent	10
3.1.2	Feed Forward Neural Network	11
3.2	Wisconsin Breast Cancer Data	14
3.2.1	Feed Forward Neural Network	14
3.2.2	Logistic Regression	14
4	Results	14
5	Discussion	14
6	Conclusion	14
A	Appendix	14

1 Introduction

In the modern world, digital data has become one of the most valuable commodities there is. Not because of scarcity, like most other valuables, but rather the exact opposite; the vast abundance of data available makes being able to understand trends and patterns in it extremely valuable for companies looking for profit. However, data is complex, having many features, and understanding how one affect another is impossible with purely human analysis. Luckily, there exists statistical methods that let us find the deeper connections, make models and even predict outcomes. In this paper we wish to study some stochastic methods, and look at their limitations and strengths.

First we will study a bi-variate continuous function known as the Franke function. We will use both stochastic gradient decent and a feed-forward deep neural network, with back propagation. Then we can also compare results with those obtained in a previous paper, using the deterministic methods ordinary least squares and ridge regression. Note that all methods used in this report is briefly covered in the theory section.

Next we will embark on a classification problem, using measurments of tumors in breast tissue to predict wether they are benign or malignant. Here we will again use a feed-forward deep neural network, along with logistic regression.

We will on no way answer all questions linked to the aforementioned methods. Such that anyone can reproduce or continue our studies, we list all the code, results and instructions on running the code in our GitHub repository¹.

2 Theory

In the theory-section we aim to give a brief explanation of the main concepts and terminology used in this report. For a more in-depth explanation we recommend reading the appropriate sections in [3], which has been of great inspiration and help for us throughout the project.

In general, we have a dataset \mathbf{x} , where each point \mathbf{x}_i takes a value y_i , for which we want to make a model β , such that for a new data point $\mathbf{x}_k \notin \mathbf{x}$ we can make a prediction for the value y_k . The model β is a vector, where each element is a parameter of our model, such that β is sometimes called the parameters. For gradient decent, we have to chose what shape the model should be, as was done for linear regression in [4], while the neural network makes its own model.

2.1 Gradient Decent

In this section we cover gradient decent and different variations of it. More specifically we describe gradient decent (GD), stochastic gradient decent (SGD) and adding momentum to the aforementioned methods. All gradient decent methods start with an initial guess for what the model β should be, and iteratively updates the guess by training on the dataset, either until it reaches a minimum, or a certain number of iterations have been performed.

¹<https://github.com/sigurdru/FYS-STK4155/tree/main/project2>

2.1.1 Ordinary Gradient Decent

Gradient decent methods is often used to minimize the so-called cost/loss-function, which tells us how good our model at predicting the dataset is (more on this in section 2.3.3). For now, we use a general cost function $C(\beta)$ for a given model β , which can be expressed as the sum over the cost function for each datapoint \mathbf{x}_i , as such:

$$C(\beta) = \sum_{i=1}^n c_i(\mathbf{x}_i, \beta), \quad (2.1)$$

where n denotes the number of datapoints. The gradient with respect to the parameters β , which represent the direction of optimal minimization of the cost function, is then defined as

$$\nabla_{\beta} C(\beta) = \sum_{i=1}^n \nabla_{\beta} c(\mathbf{x}_i, \beta). \quad (2.2)$$

The algorithm for GD is then:

$$\begin{aligned} \mathbf{v}_t &= \eta \nabla_{\beta} C(\beta_t) \\ \beta_{t+1} &= \beta_t - \mathbf{v}_t, \end{aligned} \quad (2.3)$$

where η is what we call the learning rate, representing the step-length to move in the optimal direction. This algorithms iteratively finds a new β_{t+1} which (ideally) decreases the cost function. This is of course not always the case, and depends on the value of η .

For a model with p parameters, the cost-function is the surface of a p -dimensional hypersurface, and minimizing this can lead to several problems. For example, if η is too big, the cost-function can diverge and never find a minimum of the hypersurface, while if η is too small we will need too many iterations to reach a minimum in reasonable time. One method of avoiding the cost-function diverging, is using a dynamic learning schedule, where the learning rate η decreases during training. Our model then makes larger steps in the beginning, and then smaller and smaller, such that we should be able to converge to a minimum, and not making too big steps, circling around it.

An additional problem is that the hypersurface is not a smooth terrain with a single minimum. Our model can potentially move down into a local minimum, which can be close to the level of the global minimum, or far worse than it. When our model converges, we have no way of knowing if we have found the optimal, global minium, or are stuck in one of the many local minima, with no way of getting out.

2.1.2 Stochastic Gradient Decent

With large datasets, a large number of computations is needed when calculating the gradient. It takes a lot of time, and the model is only updated once per iteration, making improvement slow. Stochastic Gradient Decent. SGD, combats this by approximating the total gradient (2.2). This is done by performing gradient decent on a subset of the data, called a minibatch. With n still denoting the total number of datapoints, we will have

$N_B = n/M$ minibatches, where M is the size of each minibatch. The minibatches are denoted by B_k . Thus our approximated gradient for a single minibatch B_k is defined as

$$\nabla_{\beta} C^{MB}(\beta) \equiv \sum_{i \in B_k}^M \nabla_{\beta} c(\mathbf{x}_i, \beta). \quad (2.4)$$

Then the aim is to use this approximated gradient, for all N_B minibatches, to update the parameters β , at every step k . Doing this for all N_B minibatches, are what we refer to as an epoch. The SGD algorithm then becomes very similar to (2.3), however with an approximated gradient.

$$\begin{aligned} \mathbf{v}_t &= \eta \nabla_{\beta} C^{MB}(\beta_t) \\ \beta_{t+1} &= \beta_t - \mathbf{v}_t \end{aligned} \quad (2.5)$$

Choosing a smaller M gives a worse approximation of the full gradient, though at the same time the model will have more chances to move in the correct direction. Further, by choosing M such that N_B is neither too small or large, our program will run faster. This not only speeds up our program, it also helps prevent getting stuck in local minima because of the stochastic nature. The dataset is shuffled after each epoch, creating new minibatches such that we never use the same one twice.

2.1.3 Adding Momentum

These methods can still be optimized further by adding momentum. This is done by adding a term to the parameter \mathbf{v}_t in equations (2.3) and (2.5). This so-called mass term, simulates the gradient having momentum, such that every update of β is a running average.

$$\mathbf{v}_t = \eta \nabla_{\beta_t} C(\beta) \rightarrow \mathbf{v}_t = \gamma \mathbf{v}_{t-1} + \eta \nabla_{\beta} C(\beta_t). \quad (2.6)$$

Here $\gamma \in [0, 1]$ is what we could refer to as the mass, and is a free parameter. One of the benefits is for example that this lets us move faster in regions where the gradient is small.

2.2 Logistic Regression

Gradient decent is used when we have a continuous output, like fitting a function to data. Logistic Regression is used for classification problems, meaning that we want to predict discrete outputs, for instance true or false, given a set of information about a subject. This is the case for binary classification, where we only have one output. It is also possible to use multi-class classification, where one has several outputs, each representing the probability of the class being given by that output.¹ In our case we want a model that takes in some values \mathbf{x}_i and spits out zero or one. These values should correspond to the actual classification $y_i \in \{0, 1\}$ corresponding to true or false respectively. Lets define our model as

$$\sigma(s_i), \quad \text{where } s_i = \mathbf{x}_i^T \mathbf{w} + b_0 \equiv \mathbf{X}_i^T \mathbf{W}, \quad (2.7)$$

¹ *kanskje flytt denne setningen til senere?*

where \mathbf{x}_i^T is our input data, \mathbf{w} and b_0 are parameters in the model. As a shorthand we also defined $\mathbf{W} = (b_0, \mathbf{w})$ and $\mathbf{X}_i = (1, \mathbf{x}_i)$. We also have σ which is some soft classifier that maps our output between zero and one (i.e. the Sigmoid (2.11)). The reason why we want a soft classifier and not a hard one (like $\sigma = 1$ if $s \geq 0$ and 0 otherwise), is because then we can interpret the output as a probability, quantifying how certain the model is in its prediction. Here we also need a cost function to minimize. It is common to choose the cross entropy, which we derive in 2.3.3. However we will just use it without derivation for now. The cross entropy for this model is given as

$$C(\mathbf{W}) = \sum_{i=1}^n -y_i \log \sigma(\mathbf{X}_i^T \mathbf{W}) - (1 - y_i) \log [1 - \sigma(\mathbf{X}_i^T \mathbf{W})], \quad (2.8)$$

where n are the number of samples we want to classify, and y_i the true classification. Now with a cost function and model in hand we are ready to minimize the cost function in order to find the optimal parameters for the model. We have a convex cost function, therefore a minimization leads to

$$\nabla_{\mathbf{W}} C(\mathbf{W}) = \sum_{i=1}^n [\sigma(\mathbf{X}_i^T \mathbf{W}) - y_i] \mathbf{X}_i. \quad (2.9)$$

Thus the only thing left to do is perform an algorithm similar to 2.5, where the parameters to update are \mathbf{W} .

2.3 Feed-Forward Deep Neural Networks

Neural networks are neural-inspired nonlinear models, which are taught by a way of learning. We will in this section explain what we mean by non-linearity, the basic architecture of a neural network and how the network learns. In this paper, only supervised learning will be studied, where we train our model on fully labeled data, as opposed to unsupervised learning, where the model first is shown unlabeled data.

2.3.1 Architecture of Neural Networks

The structure we are going to use in this report is similar to that in figure 1. The gray circles are what we refer to as nodes. For now we just need to know that they hold some numerical value. One initializes the network by giving the nodes in the input layer numerical values. These values would correspond to some actual physical property, for example brightness of pixels in a picture. Then, each node in the input layer is connected to each node in the hidden layer h_1 . In figure 1 we have three such hidden layers, where each node in one layer is connected to every node in the next layer. Now the nodes are connected through what we will refer to as weights, biases and activation functions (more on that later). The connections are what assigns the numerical value of the nodes in the next layer. Lastly we have the output layer, which outputs values dependent on the problem. If we have a classification situation, where we for example wanted to classify the type of animal in different pictures, then one node could correspond to a lion, next to a zebra and so fourth. By this we would know what animal the network *thinks* is in the picture by looking at what neuron has the highest numerical value.

The structure for the neural net we use in this report is similar to that in figure 1. The gray circles are what we refer to as nodes, organized into layers. They hold some numerical value. The first and last layer are the input and output layer, respectively. The rest are called hidden layers, denoted h_i . In figure 1 three hidden layers are shown. Between each layer is a set of weights, connecting each node in the preceding layer to each node in the succeeding one. One initializes the network by giving the nodes in the input layer numerical values. These values would correspond to some actual physical property, for example brightness of pixels in a picture. Then, the values are fed forward, the values of the nodes in the next layer being the weighted sum of the values in the previous, plus a bias b . The values are *activated* by an activation function, before similarly being fed forward to the next layer, until reaching the output layer. In the case of regression, the output layer is a single node, giving the functional value. In the case of classification, for example classifying the type of animal in pictures, one node could correspond to a lion, next to a zebra and so fourth. By this we would know what animal the network predicts is pictured by finding the neuron has the highest numerical value.




Figure 1. Basic outline of a neural network. It displays the different layers (input, hidden and output), nodes (gray circles) and the connection between the nodes (black lines).

(source: https://www.researchgate.net/figure/Artificial-neural-network-architecture-ANN-i-h-1-h-2-h-n-o_fig1_321259051)

We mentioned that the different nodes are connected through weights, biases and activation functions. Looking at figure 1, a neuron j in layer h_1 is connected to n input neurons, denoted by black lines. Each input neuron has a numerical value defined by the problem. The value neuron j in h_1 then gets is defined as

$$a_{h_1,j} = \sigma(x_1w_1 + x_2w_2 + \cdots + x_nw_n + b_j), \quad (2.10)$$

where x_i are the values of neuron i in the input layer, w_i are the weights between neurons i and j , b is what we refer to as the bias and σ is the activation function. Every neuron is connected like this, with different weights and biases. Initially the network will make random predictions. Through training, the weights and biases are updated using the back-propagation algorithm, described in section 2.3.4, until the predictions become extremely accurate.

2.3.2 Activation Functions

The activation functions are where the non-linearity of the neural nets comes in, because they are non-linear. Now there are many such functions, in our project we have implemented the 3 displayed in figure 2. The exact functions are as follows

$$\text{Sigmoid: } \sigma(x) = \frac{1}{1 + e^{-x}}, \quad (2.11)$$

$$\text{ReLU: } \sigma(x) = \max(0, x), \quad (2.12)$$

$$\text{Leaky ReLU: } \sigma(x) = \begin{cases} \alpha x, & \text{if } x \leq 0 \\ x, & \text{otherwise.} \end{cases} \quad (2.13)$$

For Leaky ReLU, α is some parameter for which we use $\alpha = 0.01$.




Figure 2. Some activation functions, Sigmoid, ReLU and Leaky ReLU, respectively.

All neurons in all hidden layers are activated by the same activation function. The input layer is just the input, so it is not activated. In the case of regression, neither the output layer is activated. Then the function value is just the weighted sum from the last hidden layer. In the case of classification, we want the values in the output nodes to indicate the probabilities. Thus, we use an activation function that considers the value in the other nodes in the other output nodes. The softmax function does this.

$$\sigma(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}},$$

where K is the number of classes, or output nodes.

2.3.3 Cost Function and Regularization

Before one can start training the data, we must have a cost function. This will tell us how well or poorly our network is performing, and is what we want to minimize when we train the network. For continuous data it is common to use mean square error (MSE) as the cost function. Which is just the difference between desired output ($\hat{\mathbf{x}}$) and actual output (\mathbf{x}), squared, divided by the number of datapoints n (see equation (2.14))

$$C(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n (\hat{\mathbf{x}}_i - \mathbf{x}_i)^2. \quad (2.14)$$

One can also implement regularization, which helps prevent overfitting in the network. Common ones are L_1 and L_2 which is done by adding a regularization term to the end of the cost function (see equations (2.15) and (2.16)).

$$L_1 : C(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n (\hat{\mathbf{x}}_i - \mathbf{x}_i)^2 + \lambda \sum_j |\mathbf{w}_j| \quad (2.15)$$

$$L_2 : C(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n (\hat{\mathbf{x}}_i - \mathbf{x}_i)^2 + \lambda \sum_j \mathbf{w}_j^2 \quad (2.16)$$

²Here λ is some regularization parameter and \mathbf{w} are weights.

² RIK-TIG?

We mentioned when talking about logistic regression (section 2.2), that for classification scenarios one will often use cross-entropy as the cost function. Continuing with the model defined by (2.7), we want to find an appropriate cost function. We define the probability of an outcome y_i given parameters \mathbf{X}_i and \mathbf{W} as

$$P(y_i = 1 | \mathbf{X}_i, \mathbf{W}) = \frac{1}{1 + \exp(-\mathbf{X}_i^T \mathbf{W})}, \quad (2.17)$$

$$P(y_i = 0 | \mathbf{X}_i, \mathbf{W}) = 1 - P(y_i = 1 | \mathbf{X}_i, \mathbf{W}) \quad (2.18)$$

We can then map these probabilities to our soft classifier $\sigma(s_i)$

$$P(y_i = 1) = \sigma(s_i) = \sigma(\mathbf{X}_i^T \mathbf{W}). \quad (2.19)$$

Now we can define the cost function using Maximum Likelihood Estimation (MLE), which states that we should choose parameters that maximize the probability of our given data. Consider the dataset $\mathcal{D}\{(y_i, \mathbf{x}_i)\}$, where we remind that \mathbf{x}_i are the input parameters. Then the probability of our dataset given \mathbf{W} is

$$P(\mathcal{D} | \mathbf{W}) = \prod_{i=1}^n \left[\sigma(\mathbf{X}_i^T \mathbf{W}) \right]^{y_i} \left[1 - \sigma(\mathbf{X}_i^T \mathbf{W}) \right]^{(1-y_i)}. \quad (2.20)$$

Again we remind that n are the number datapoints we want to classify. This expression is difficult to work with, thus we take the logarithm.

$$l(\mathbf{W}) = \log(P(\mathcal{D} | \mathbf{W})) = \sum_{i=1}^n y_i \log(\sigma(\mathbf{X}_i^T \mathbf{W})) + (1 - y_i) \log(1 - \sigma(\mathbf{X}_i^T \mathbf{W})) \quad (2.21)$$

MLE entails finding the \mathbf{W} that maximizes $l(\mathbf{W})$, or more commonly, minimizes $-l(\mathbf{W})$. Thus our cost function becomes

$$C(\mathbf{W}) = -l(\mathbf{W}) = \sum_{i=1}^n -y_i \log \sigma(\mathbf{X}_i^T \mathbf{W}) - (1 - y_i) \log [1 - \sigma(\mathbf{X}_i^T \mathbf{W})], \quad (2.22)$$

Which is exactly what we wrote down in equation (2.8).

2.3.4 The Backpropagation Algorithm

With a desired cost function we are ready to train the neural network. This is done by the backpropagation algorithm. The method entails finding the derivative of the cost function, with respect to all parameters. When we have a neural network, we have thousands of parameters which can be tuned (weights and biases), meaning that we have to approximate the derivative somehow. The backpropagation algorithm does just that, by exploiting the layered structure displayed in figure 1.

Before we can embark on deriving the algorithm we will introduce some notation. We assume L total layers while $l = 1, \dots, L$ indexes which one. Next we need to index the weights, nodes and biases. Let w_{jk}^l be the weight connecting k -th neuron in layer $l - 1$ and j -th neuron in layer l . The index order in j and k are such that we can do matrix multiplication with index notation later down the road. We also let b_j^l be the j -th neuron bias in layer l . Thus the activation of the j -th neuron in layer l (a_j^l) becomes

$$a_j^l = \sigma \left(\sum_k w_{jk}^l a_k^{l-1} + b_j^l \right) = \sigma(z_j^l), \quad z_j^l \equiv \sum_k w_{jk}^l a_k^{l-1} + b_j^l. \quad (2.23)$$

Here σ is an activation function.

Now the cost function will depend directly on the activation of the output layer (a_j^L). However the activation of the output layer depends on the previous layers, meaning that the cost function depends indirectly on all the previous layers. Lets define the error Δ_j^L of the j -th neuron in layer L , as the change in cost function with respect to z_j^L .

$$\Delta_j^L \equiv \frac{\partial C}{\partial z_j^L} \quad (2.24)$$

We can similarly define the error of neuron j in layer l (Δ_j^l), as the change in the cost function with respect to z_j^l ,

$$\Delta_j^l \equiv \frac{\partial C}{\partial z_j^l} = \frac{\partial C}{\partial a_j^l} \frac{\partial a_j^l}{\partial z_j^l} = \frac{\partial C}{\partial a_j^l} \frac{d\sigma(z_j^l)}{dz_j^l}. \quad (2.25)$$

In the next few lines we are going to derive several equations needed for the algorithm, it will be apparent why after we have found them. Notice that (2.25) also can be written as

$$\Delta_j^l = \frac{\partial C}{\partial z_j^l} = \frac{\partial C}{\partial b_j^l} \frac{\partial b_j^l}{\partial z_j^l} = \frac{\partial C}{\partial b_j^l}. \quad (2.26)$$

Because $\partial b_j^l / \partial z_j^l = 1$ from (2.23). Again using the chain rule we can rewrite (2.25)

$$\begin{aligned} \Delta_j^l &= \frac{\partial C}{\partial z_j^l} = \sum_k \frac{\partial E}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} = \sum_k \Delta_k^{l+1} \frac{\partial z_k^{l+1}}{\partial z_j^l} \\ &= \left(\sum_k \Delta_k^{l+1} w_{kj}^{l+1} \right) \frac{d\sigma(z_j^l)}{dz_j^l}. \end{aligned} \quad (2.27)$$

To find the last equation, we differentiate the cost function with respect to the weight w_{jk}^l

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l} = \Delta l a_k^{l-1}. \quad (2.28)$$

Now why have we done all this work, well because the equations (2.25), (2.26), (2.27) and (2.28) define what we call the backpropagation algorithm. Then, what exactly is the algorithm? It entails six steps:

- 1 Activation:** First activate the neurons in the activation layer (a_j^1) with desired data.
- 2 Feedforward:** Activate the nodes in following layers, this is done by equation (2.23).
- 3 Error at layer L :** Calculate the error at the last layer using (2.25).
- 4 Backpropagate error:** With (2.27) we can calculate the error, iterating backwards in the network.
- 5 Calculate gradient:** Find the gradient by using equations (2.26) and (2.28).
- 6 Update parameters:** Update the parameters similarly to (2.5), however β_t are our weights and biases in this case.³

³ RIK-TIG?

The expression for updating the weights and biases are

$$w_{jk}^l \leftarrow w_{jk}^l - \eta \Delta_j^l a_k^{l-1} \quad (2.29)$$

$$b_j^l \leftarrow b_j^l - \eta \Delta_j^l \quad (2.30)$$

2.3.5 Initialization of weights

We mentioned earlier than when the network is created, it has weights and biases between the layers. These need to be initialized in some way. The biases are simple to initialize, as they are a single number for every node. These are initialized as a small, non-zero value b_0 , which we choose as $b_0 = 0.01$.

Before 2006, most neural networks were performing quite badly on most tasks, as they did not learn during training. One of the (several) reasons were due to bad initialization of weights. A common way of doing this was using the standard normal distribution $W_{i,j} \sim \mathcal{N}(0, 1)$. The problem with this is that it does not consider the size of the layers. In 2010, it was shown that when using sigmoid as the activation function, Xavier-initialization give better results [1]. This is given as $W_{i,j} \sim \mathcal{U}\left(-\frac{1}{\sqrt{n}}, \frac{1}{\sqrt{n}}\right)$, where \mathcal{U} is the uniform distribution, and n is the number of nodes in the preceding layer.

In 2015, He-initialization was shown to work well with ReLU and Leaky ReLU [2]. Here, the weights are initialized using the normal distribution, but with a variance given by $v = 2/(1 + \alpha^2)n$, where again n is the number of nodes in the preceding layer, and α is the parameter of the Leaky ReLU-function. For ReLU, this is 0. These initializations only consider the number of nodes in the preceding layer, though normalized initializations considering also the number of nodes in the succeeding layer could yield better results.

We use the initializations as given here, for the given activation-function, and will not study the particular effects of this in depth.

3 Methods

As we mentioned in the introduction, we wish to study different ways of fitting two types of datasets. The first which we can classify as *continuous* is the Franke Function (3.1),

$$f(x, y) = \frac{3}{4} \exp\left(-\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4}\right) + \frac{3}{4} \exp\left(-\frac{(9x+1)^2}{49} - \frac{(9y+1)^2}{10}\right) + \frac{1}{2} \exp\left(-\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4}\right) - \frac{1}{5} \exp(-(9x-4)^2 - (9y-7)^2). \quad (3.1)$$

We will use both stochastic gradient decent and feed forward neural network to try and fit the data. Both methods are covered in the theory sections 2.1 and 2.3 respectively.

Next we will embark on an classification problem. Namely classifying if breast tissue is malignant or benign, by studying the data provided by Wisconsin breast cancer data². We will again use a feed forward neural network, and logistic regression. The latter is covered in the theory section 2.2.

3.1 Franke Function

In [4], we already studied the Franke function using linear regression, specifically OLS and OLS with an L2 and L1 parameter λ , so-called Ridge and Lasso regression. The results from these methods will form the basis for comparing our results using SGD and neural networks. The rapport, along with the code can be found at our GitHub³. In order to have comparable results, we will use the same parameters for the data. Only the methods will be different. In that project we initialized the data with $N = 30 \times 30$ uniformly distributed datapoints in x - and y -direction. To simulate it being real data, we also added normally distributed noise with mean zero and standard deviation 0.2: $\epsilon \sim \mathcal{N}(0, 0.2)$. We also split the input and target data in the same way as before, using 80 % of the data for training and 20 % of the data for testing. The two splitted data sets are then scaled by subtracting the mean of the relevant training data.

3.1.1 Stochastic Gradient Descent

As in [4], we have to choose a model to fit the data to, when using SGD. The simplest is a bi-variate polynomial of degree P , such that our model will have p features. This is the design matrix X used in the previous project. Having obtained good results for OLS using $P = 6$, we use the same polynomial degree for SGD. Writing our own code for implementing SGD, we will analyze the results with the MSE (equation (2.14)) as our cost function for various parameters. We include the L_2 regularization term, equation (2.16), in the gradient of the cost function, which gives the gradient corresponding to Ridge regression. The expression for the stochastic gradient is given in equation (3.2)

²<https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

³<https://github.com/sigurdru/FYS-STK4155/tree/main/project1>

$$\nabla_{\beta} C^{\text{MB}}(\beta) = 2X^T(X\beta - \mathbf{z})/M + 2\lambda\beta \quad (3.2)$$

where OLS regression is obtained by setting $\lambda = 0$. There are multiple parameters to consider for stochastic gradient descent, and the ones we will study are the choice of learning rates, η , number of epochs, N_e , number of minibatches, N_B , regularization parameter, λ and the momentum parameter γ .

We begin by studying the MSE for different values of η as a function of $N_e \in [0, 150]$, setting $\lambda = \gamma = 0$ and $N_B = 20$, i.e. minibatches of size 36 for the 720 training data.

⁴ We choose 21 evenly spaced values of $\eta \in [0.01, 0.5]$. This will simplify our analysis, as studying all the parameters at once are a tedious process. We then choose one of the favorable learning rates to study the MSE for 9 different minibatch sizes between 720 and 24, as a function of epochs, using $N_e \in [0, 150]$ once again. Using $N_B = 20$ minibatches, we then study the MSE after $N_e = 150$ epochs as a function of η and λ to study the effect of regularization. For this we use 11 evenly distributed values of $\eta \in [0.1, 0.5]$ and 11 logarithmically distributed values of $\lambda \in [10^{-5}, 1]$. Finally, we study the MSE over 100 epochs for 20 evenly spaced values of $\gamma \in [0, 0.95]$ ⁵.

After these parameters are studied, we will test an implementation of scaling the learning rate by reducing it from an initial value after each epoch. One possible advantage of this is that convergence is fast initially, when the errors are large, but as we approach a minima the learning rates are reduced, preventing us from overshooting. The algorithm we use for scaling the learning rate is given in equation (3.3), where t denotes the epoch

$$\eta_t = \eta_0 \cdot \left(1 - \frac{t}{N_e}\right) \quad (3.3)$$

Parameter for learning schedule ⁶.

After the SGD has been studied and appropriate parameters have been estimated, we will plot the resulting fit of the gradient descent and compare it to the Franke function as an indication of the result. Although the fit was performed on a data set with noise $\epsilon \sim \mathcal{N}(0, 0.2)$ we will use $\epsilon \sim \mathcal{N}(0, 0.05)$ when plotting the prediction in order to see the details of the plot more clearly.

3.1.2 Feed Forward Neural Network

When using the neural network to fit the Franke function, we use a lot of the same methods as for SGD. A key difference is that instead of iteratively updating a model β , we now train a network of several layers, each with many nodes. One of the results of this is that we do not have to choose the shape the model will take.

Instead of giving our network the design matrix X for a certain polynomial degree P , we can pass it only the collection of points $[(x_i, y_i)]$, and let the network adjust the weights and biases accordingly. Since the Franke function is an exponential function we know that

⁴ skal vi kommentere her at best eta påvirkes av batch size, eller skal det være i diskusjonen?
⁵ Burde man kanskje lagd tabell med parametrene?

⁶ TO BE DONE

it can be approximated as a higher order polynomial, so by using the design matrix as an input we exploit this property such that the network converges faster. Having already fitted the Franke function with a design matrix with linear regression and SGD, we now choose the x_i and y_i values only. Not providing the network with any initial information has the advantage that we get a more rigorous test of the network’s performance, since we ensure that the result is not directly reliant on the information in question. Another important motivation for this is that if we were to fit some other data, e.g. terrain data, we may not have any a priori information regarding the input data. Omitting the design matrix when we train our neural network will thus yield a final model capable of fitting various types of data.

Since we are dealing with a regression problem and we’re fitting a continuous function, a natural choice of the cost function is the MSE. For the neural network we will not compute the total MSE of the output layer as we have previously done, but the individual MSE of each output node. This takes into account the error at each individual output neuron when we update the weights with backpropagation.

«««< HEAD The first thing we will do is to study the evolution of the output layers’ MSE for different values of η as we train the network. This will give an estimate of reasonable learning rates to use for training. We choose two hidden layers in the neural network with ten nodes each where we use the sigmoid activation function from equation (2.11). The first thing we must consider is how we initialize the weights and biases for the different layers. We initialize the weights randomly using a normal distribution of $\mathcal{N}(0, 1)$. The biases are a single non-zero number for each node, and we choose $b_0 = 0.01$ for all nodes initially. For the output layer we don’t use any activation function, since the output of the Neural network should be the prediction of the Franke function⁴. For the initial analysis we will not include regularization or a momentum parameter.

Having initialized the neural network we are now going to train it. Each training iteration begins by randomly shuffling the data and dividing them into minibatches, just as we did when we performed the SGD analysis. For each minibatch we begin by using the input data to update each layer of the network until we reach the output layer. Then we use the backpropagation algorithm to update the weights and biases at every node. **ER AVS-NITTET OVER UNØDVENDIG?** With no activation function in the output layer, the error of this layer is given by equation (2.24) i.e. the derivative of the MSE. Iterating backwards we get the error in the previous layers by using equation (2.27). The weights and biases at each layer are then updated with equations (2.29) and (2.30) respectively. Doing this for all minibatches we complete one epoch.

For the first simulation we choose 21 linearly distributed values of $\eta \in [10^{-3}, 0.9]$. We will then repeat this analysis using 21 new η values on an interval where the resulting MSE was low.

FUCK VI MANGLER R2.

⁴This is a reasonable choice considering that the maximum of equation (3.1) exceeds 1, and the minimum is negative when noise is added

We now study regularization parameters. HÅKON OLAV

So far, we have only studied the statistical accuracy of our neural network. As a sanity check we will use the optimal parameters obtained in order to plot the resulting prediction of the neural network, using $\epsilon \sim \mathcal{N}(0, 0.05)$ just as we did for SGD. Some of the parameters we use may not be ideal when the noise is reduced, but the purpose of this test is simply to see whether our model behaves in a desired way, so reducing the noise makes the comparison easier.

Having only studied different learning rates and regularization parameters of the neural network, we are now ready to test different configurations of the network itself. The first modification we will make is testing the neural network using different activation functions for the hidden layers. Combined with this, we will apply different initialization methods of the weights and biases ⁷.

===== We choose two hidden layers in the neural network with ten nodes each for the Franke function. For the hidden layers we use the sigmoid activation function from equation (2.11). To train our network we first initialize the weights and biases for the different layers. We initialize the weights randomly using a normal distribution of $\mathcal{N}(0, 1)$. The biases are a single non-zero number for each node, and we choose $b_0 = 0.01$ for all nodes initially. For the output layer we don't use any activation function, meaning that the output of the Neural network should be the prediction of the Franke function.

⁷ *biases
har vi vel
ikke gjort*

Having initialized the neural network we are now going to train it. Each training iteration begins by randomly shuffling the data and dividing them into minibatches, just as we did when we performed the SGD analysis. For each minibatch we begin by using the input data to update each layer of the network until we reach the output layer. Then we use the backpropagation algorithm to update the weights and biases at every node. With no activation function in the output layer, the error of this layer is given by equation (2.24) i.e. the derivative of the MSE. Iterating backwards we get the error in the previous layers by using equation (2.27). The weights and biases at each layer are then updated with equations (2.29) and (2.30) respectively. Doing this for all minibatches we complete one epoch. We initially train our network »»»> 670ff5071b1da8764c091f5e90b02cff9b5714d8

3.2 Wisconsin Breast Cancer Data

3.2.1 Feed Forward Neural Network

3.2.2 Logistic Regression

4 Results

5 Discussion

6 Conclusion

A Appendix

References

- [1] Xavier Glorot and Y. Bengio. Understanding the difficulty of training deep feedforward neural networks. *Journal of Machine Learning Research - Proceedings Track*, 9:249–256, 01 2010.
- [2] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. *IEEE International Conference on Computer Vision (ICCV 2015)*, 1502, 02 2015.
- [3] Pankaj Mehta, Marin Bukov, Ching-Hao Wang, Alexandre G.R. Day, Clint Richardson, Charles K. Fisher, and David J. Schwab. A high-bias, low-variance introduction to machine learning for physicists. *Physics Reports*, 810:1–124, May 2019.
- [4] Håkon Olav Torvik, Vetle Vikenes, and Sigurd Sørle Rustad. Analysis of regression and resampling methods. *University of Oslo*, October 2021.