

Machine Learning: Using regression and neural networks to fit continuous functions and classify data



Håkon Olav Torvik, Vetle Vikenes & Sigurd Sørli Rustad

FYS-STK4155 – Applied Data Analysis and Machine Learning

Autumn 2021

Department of Physics

University of Oslo

November 15, 2021

ABSTRACT: Abstract coming soon.

Contents

1	Introduction	1
2	Theory	1
2.1	Gradient Decent	1
2.1.1	Ordinary Gradient Decent	2
2.1.2	Stochastic Gradient Decent	2
2.1.3	Adding Momentum	3
2.2	Logistic Regression	3
2.3	Feed-Forward Deep Neural Networks	4
2.3.1	Architecture of Neural Networks	4
2.3.2	Activation Functions	5
2.3.3	Cost Function and Regularization	6
2.3.4	The Backpropagation Algorithm	7
2.3.5	Initialization of weights	9
3	Methods	9
3.1	Franke Function	10
3.1.1	Stochastic Gradient Descent	10
3.1.2	Feed Forward Neural Network	12
3.2	Wisconsin Breast Cancer Data	13
3.2.1	Feed Forward Neural Network	13
3.2.2	Logistic Regression	13
4	Results	14
4.1	Franke Function	14
4.1.1	Stochastic Gradient Descent	14
4.1.2	Feed Forward Neural Network	19
4.2	Wisconsin Breast Cancer Data	23
4.2.1	Feed Forward Neural Network	23
4.2.2	Logistic Regression	26
5	Discussion	29
6	Conclusion	29

1 Introduction

In the modern world, digital data has become one of the most valuable commodities there is. Not because of scarcity, like most other valuables, but rather the exact opposite; the vast abundance of data available makes being able to understand trends and patterns in it extremely valuable for companies looking for profit. However, data is complex, having many features, and understanding how one affect another can be impossible with human analysis alone. Luckily, there exists statistical methods that let us find the deeper connections, make models and even predict outcomes. In this paper we wish to study some stochastic methods, and look at their limitations and strengths.

First we will study a bi-variate continuous function known as the Franke function. We will use both stochastic gradient decent and a feed-forward deep neural network, with back propagation. Then we can also compare results with those obtained in a previous paper, using the deterministic methods ordinary least squares and ridge regression. Note that all methods used in this report is briefly covered in the theory section.

Next we will embark on a classification problem, using measurments of tumors in breast tissue to predict wether they are benign or malignant. Here we will again use a feed-forward deep neural network, along with logistic regression.

We will in no way answer all questions linked to the aforementioned methods. So that anyone can reproduce or continue our studies, we list all the code, results and instructions on running the code in our GitHub repository¹.

2 Theory

In the theory-section we aim to give a brief explanation of the main concepts and terminology used in this report. For a more in-depth explanation we recommend reading the appropriate sections in [3], which has been of great inspiration and help for us throughout the project.

In general, we have a dataset \mathbf{x} , where each point \mathbf{x}_i takes a value y_i , for which we want to make a model β , such that for a new data point $\mathbf{x}_k \notin \mathbf{x}$ we can make a prediction for the value y_k . The model β is a vector/matrix, where each element is a parameter of our model, such that β is sometimes called the parameters. For gradient decent, we have to chose what shape the model should be, as was done for linear regression in [4], while the neural network makes its own model.

2.1 Gradient Decent

In this section we cover gradient decent and different variations of it. More specifically we describe gradient decent (GD), stochastic gradient decent (SGD) and adding momentum to the aforementioned methods. All gradient decent methods start with an initial guess for what the model β should be, and iteratively updates the guess by training on the dataset, either until it reaches a minimum, or a certain number of iterations have been performed.

¹<https://github.com/sigurdru/FYS-STK4155/tree/main/project2>

2.1.1 Ordinary Gradient Decent

Gradient decent methods is often used to minimize the so-called cost/loss-function, which tells us how good our model at predicting the dataset is (more on this in section 2.3.3). For now, we use a general cost function $C(\beta)$ for a given model β , which can be expressed as the sum over the cost function for each datapoint \mathbf{x}_i , as such:

$$C(\beta) = \sum_{i=1}^n c_i(\mathbf{x}_i, \beta), \quad (2.1)$$

where n denotes the number of datapoints. The gradient with respect to the parameters β , which represent the direction of optimal minimization of the cost function, is then defined as

$$\nabla_{\beta} C(\beta) = \sum_{i=1}^n \nabla_{\beta} c(\mathbf{x}_i, \beta). \quad (2.2)$$

The algorithm for GD is then:

$$\begin{aligned} \mathbf{v}_t &= \eta \nabla_{\beta} C(\beta_t) \\ \beta_{t+1} &= \beta_t - \mathbf{v}_t, \end{aligned} \quad (2.3)$$

where η is what we call the learning rate, representing the step-length to move in the optimal direction. This algorithms iteratively finds a new β_{t+1} which (ideally) decreases the cost function. This is of course not always the case, and depends on the value of η .

For a model with p parameters, the cost-function is the surface of a p -dimensional hypersurface, and minimizing this can lead to several problems. For example, if η is too big, the cost-function can diverge and never find a minimum of the hypersurface, while if η is too small we will need too many iterations to reach a minimum in reasonable time. One method of avoiding the cost-function diverging, is using a dynamic learning schedule, where the learning rate η decreases during training. The model then makes larger steps in the beginning, and then smaller and smaller, such that we should be able to converge to a minimum, and not making too big steps, circling around it.

An additional problem is that the hypersurface is not a smooth terrain with a single minimum. Our model can potentially move down into a local minimum, which can be close to the level of the global minimum, or far worse than it. When our model converges, we have no way of knowing if we have found the optimal, global minium, or are stuck in one of the many local minima, with no way of getting out.

2.1.2 Stochastic Gradient Decent

With large datasets, a large number of computations is needed when calculating the gradient. It takes a lot of time, and the model is only updated once per iteration, making improvement slow. Stochastic Gradient Decent. SGD, combats this by approximating the total gradient (2.2). This is done by performing gradient decent on a subset of the data, called a minibatch. With n still denoting the total number of datapoints, we will have

$N_B = n/M$ minibatches, where M is the size of each minibatch. The minibatches are denoted by B_k . Thus our approximated gradient for a single minibatch B_k is defined as

$$\nabla_{\beta} C^{MB}(\beta) \equiv \sum_{i \in B_k}^M \nabla_{\beta} c(\mathbf{x}_i, \beta). \quad (2.4)$$

Then the aim is to use this approximated gradient, for all N_B minibatches, to update the parameters β , at every step k . Doing this for all N_B minibatches, are what we refer to as an epoch. The SGD algorithm then becomes very similar to (2.3), however with an approximated gradient.

$$\begin{aligned} \mathbf{v}_t &= \eta \nabla_{\beta} C^{MB}(\beta_t) \\ \beta_{t+1} &= \beta_t - \mathbf{v}_t \end{aligned} \quad (2.5)$$

Choosing a smaller M gives a worse approximation of the full gradient, though at the same time the model will have more chances to move in the correct direction. Further, by choosing M such that N_B is neither too small or large, our program will run faster. This not only speeds up our program, it also helps prevent getting stuck in local minima because of the stochastic nature. The dataset is shuffled after each epoch, creating new minibatches such that we never use the same one twice.

2.1.3 Adding Momentum

These methods can still be optimized further by adding momentum. This is done by adding a term to the parameter \mathbf{v}_t in equations (2.3) and (2.5). This so-called mass term, simulates the gradient having momentum, such that every update of β is a running average.

$$\mathbf{v}_t = \eta \nabla_{\beta_t} C(\beta) \rightarrow \mathbf{v}_t = \gamma \mathbf{v}_{t-1} + \eta \nabla_{\beta} C(\beta_t). \quad (2.6)$$

Here $\gamma \in [0, 1]$ is what we could refer to as the mass, and is a free parameter. One of the benefits is for example that this lets us move faster in regions where the gradient is small.

2.2 Logistic Regression

Logistic Regression is used for classification problems, meaning that we want to predict discrete outputs, for instance true or false, given a set of information about a subject. This is the case for binary classification, where we only have one output. It is also possible to use multi-class classification, where one has several outputs, each representing the probability of the class being given by that output. In our case we want a model that takes in some values \mathbf{x}_i and spits out zero or one. These values should correlate with the actual classification $y_i \in \{0, 1\}$ corresponding to true or false respectively. Lets define our model as

$$\sigma(s_i), \text{ where } s_i = \mathbf{x}_i^T \mathbf{w} + b_0 \equiv \mathbf{X}_i^T \mathbf{W}, \quad (2.7)$$

where \mathbf{x}_i^T is our inputdata, \mathbf{w} and b_0 are parameters in the model. As a shorthand we also defined $\mathbf{W} = (b_0, \mathbf{w})$ and $\mathbf{X}_i = (1, \mathbf{x}_i)$. We also have σ which is some soft classifier that

maps our output between zero and one (i.e. the Sigmoid (2.11)). The reason why we want a soft classifier and not a hard one (like $\sigma = 1$ if $s \geq 0$ and 0 otherwise), is because then we can interpret the output as a probability, quantifying how certain the model is in its prediction. Here we also need a cost function to minimize. It is common to choose the cross entropy, which we derive in 2.3.3. However we will just use it without derivation for now. The cross entropy for this model is given as

$$C(\mathbf{W}) = \sum_{i=1}^n -y_i \log \sigma(\mathbf{X}_i^T \mathbf{W}) - (1 - y_i) \log [1 - \sigma(\mathbf{X}_i^T \mathbf{W})], \quad (2.8)$$

where n are the number of samples we want to classify, and y_i the true classification. Now with a cost function and model in hand we are ready to minimize the cost function in order to find the optimal parameters for the model. The gradient of the cost function is given by

$$\nabla_{\mathbf{W}} C(\mathbf{W}) = \sum_{i=1}^n [\sigma(\mathbf{X}_i^T \mathbf{W}) - y_i] \mathbf{X}_i. \quad (2.9)$$

Thus the only thing left to do is perform an algorithm similar to 2.5, where the parameters to update are \mathbf{W} .

2.3 Feed-Forward Deep Neural Networks

Neural networks are neural-inspired nonlinear models, which are taught by a way of learning. We will in this section explain what we mean by non-linearity, the basic architecture of a neural network and how the network *learns*. In this paper, only supervised learning will be studied, where we train our model on fully labeled data, as opposed to unsupervised learning, where the model first is shown unlabeled data.

2.3.1 Architecture of Neural Networks

The structure of the neural net we will use in this report is similar to that in figure 1. The gray circles are what we refer to as nodes, and are organized into layers. They hold some numerical value. The first and last layer are the input and output layer, respectively. The rest are so-called hidden layers, denoted h_i . In figure 1 there are three hidden layers. Between each layer is a set of weights, connecting each node in the preceding layer to each node in the succeeding one. One initializes the network by giving the nodes in the input layer numerical values. These values would correspond to some actual physical property, for example brightness of pixels in a picture. Then, the values are fed forward, where the values of the nodes in the next layer being the weighted sum of the values in the previous, plus a bias b . The values are *activated* by an activation function, before similarly being fed forward to the next layer, until reaching the output layer. In the case of regression, the output layer is a single node, giving the functional value. In the case of classification, for example classifying the type of animal in pictures, one node could correspond to a lion, next to a zebra and so fourth. By this we would know what animal the network predicts is pictured by finding the neuron has the highest numerical value.



Figure 1. Basic outline of a neural network. It displays the different layers (input, hidden and output), nodes (gray circles) and the connection between the nodes (black lines).

(source: https://www.researchgate.net/figure/Artificial-neural-network-architecture-ANN-i-h-1-h-2-h-n-o_fig1_321259051)

We mentioned that the different nodes are connected through weights, biases and activation functions. Looking at figure 1, a neuron j in layer h_1 is connected to n input neurons, denoted by black lines. Each input neuron has a numerical value defined by the problem. The value neuron j in h_1 then gets is defined as

$$a_{h_1,j} = \sigma(x_1w_1 + x_2w_2 + \cdots + x_nw_n + b_j), \quad (2.10)$$

where x_i are the values of neuron i in the input layer, w_i are the weights between neurons i and j , b is what we refer to as the bias and σ is the activation function. Every neuron is connected like this, with different weights and biases. Initially the network will make random predictions. Through training, the weights and biases are updated using the backpropagation algorithm, described in section 2.3.4.

2.3.2 Activation Functions

The activation functions are where the non-linearity of the neural nets comes in, because they are non-linear. Now there are many such functions, in our project we have implemented the 3 displayed in figure 2. The exact functions are as follows

$$\text{Sigmoid: } \sigma(x) = \frac{1}{1 + e^{-x}}, \quad (2.11)$$

$$\text{ReLU: } \sigma(x) = \max(0, x), \quad (2.12)$$

$$\text{Leaky ReLU: } \sigma(x) = \begin{cases} \alpha x, & \text{if } x \leq 0 \\ x, & \text{otherwise.} \end{cases} \quad (2.13)$$

For Leaky ReLU, α is some parameter for which we set to $\alpha = 0.01$. It can theoretically happen that all the input to the activation function is less than 0. As we will show in section 2.3.4, the network weight and biases are updated using the derivative of the activation function. Using ReLU, all the derivatives would in this case be 0, and the neurons effectively become dead, never learning. This is the main benefit of using Leaky ReLU, as this won't happen. However, using good weight initialization, this should not be

a problem, so we do not expect the difference between the two activation functions to be very significant.




Figure 2. Some activation functions, Sigmoid, ReLU and Leaky ReLU, respectively.

All neurons in all hidden layers are activated by the same activation function. The input layer is just the input, so it is not activated. In the case of regression, the output layer is not activated. Then the predicted function value is just the weighted sum from the last hidden layer. For classification, we want the values in the output nodes to indicate probabilities. Thus, we use an activation function that considers the value in the other nodes in the other output nodes. The softmax function does this

$$\sigma(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}, \quad (2.14)$$

where K is the number of classes, output nodes.

2.3.3 Cost Function and Regularization

Before one can start training the data, we must have a cost function. This will quantify how well or poorly our network is performing, and is what we want to minimize when we train the network. For continuous data it is common to use mean square error (MSE) as the cost function. It is just the difference between desired output ($\hat{\mathbf{x}}$) and actual output (\mathbf{x}), squared, averaged over all datapoints \mathbf{x}_i , as such

$$C(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n (\hat{\mathbf{x}}_i - \mathbf{x}_i)^2. \quad (2.15)$$

As was done with linear regression, [4] regularization can help prevent overfitting. Common regularization methods are L_1 and L_2 penalties, which add a regularization term to the cost function, as such

$$L_1 : C(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n (\hat{\mathbf{x}}_i - \mathbf{x}_i)^2 + \lambda \sum_j |\mathbf{W}_j|, \quad (2.16)$$

$$L_2 : C(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n (\hat{\mathbf{x}}_i - \mathbf{x}_i)^2 + \lambda \sum_j \mathbf{W}_j^2, \quad (2.17)$$

where λ is some regularization parameter and \mathbf{W} are the weights and biases. We will in this paper study the effects of adding a L_2 -regulizer term when updating the parameters,

by adding the term $\eta\lambda\mathbf{W}_{t-1}$ to equation (2.6).

We mentioned when talking about logistic regression (section 2.2), that for classification scenarios one will often use cross-entropy as the cost function. Continuing with the model defined by (2.7), we want to find an appropriate cost function. We define the probability of an outcome y_i given parameters \mathbf{X}_i and \mathbf{W} as

$$P(y_i = 1|\mathbf{X}_i, \mathbf{W}) = \frac{1}{1 + \exp(-\mathbf{X}_i^T \mathbf{W})}, \quad (2.18)$$

$$P(y_i = 0|\mathbf{X}_i, \mathbf{W}) = 1 - P(y_i = 1|\mathbf{X}_i, \mathbf{W}) \quad (2.19)$$

We can then map these probabilities to our soft classifier $\sigma(s_i)$

$$P(y_i = 1) = \sigma(s_i) = \sigma(\mathbf{X}_i^T \mathbf{W}). \quad (2.20)$$

Now we can define the cost function using Maximum Likelihood Estimation (MLE), which states that we should choose parameters that maximize the probability of our given data. Consider the dataset $\mathcal{D}\{(y_i, \mathbf{x}_i)\}$, where we remind that \mathbf{x}_i are the input parameters. Then the probability of our dataset given \mathbf{W} is

$$P(\mathcal{D}|\mathbf{W}) = \prod_{i=1}^n [\sigma(\mathbf{X}_i^T \mathbf{W})]^{y_i} [1 - \sigma(\mathbf{X}_i^T \mathbf{W})]^{(1-y_i)}. \quad (2.21)$$

Again we remind that n are the number datapoints we want to classify. This expression is difficult to work with, thus we take the logarithm.

$$l(\mathbf{W}) = \log(P(\mathcal{D}|\mathbf{W})) = \sum_{i=1}^n y_i \log(\sigma(\mathbf{X}_i^T \mathbf{W})) + (1 - y_i) \log(1 - \sigma(\mathbf{X}_i^T \mathbf{W})) \quad (2.22)$$

MLE entails finding the \mathbf{W} that maximizes $l(\mathbf{W})$, or more commonly, minimizes $-l(\mathbf{W})$. Thus our cost function becomes

$$C(\mathbf{W}) = -l(\mathbf{W}) = \sum_{i=1}^n -y_i \log \sigma(\mathbf{X}_i^T \mathbf{W}) - (1 - y_i) \log [1 - \sigma(\mathbf{X}_i^T \mathbf{W})], \quad (2.23)$$

which is equation (2.8).

2.3.4 The Backpropagation Algorithm

With a desired cost function we are ready to train the neural network. This is done by the backpropagation algorithm. The method entails finding the derivative of the cost function, with respect to all parameters. When we have a neural network, we have thousands of parameters which can be tuned (weights and biases), meaning that we have to approximate the derivative somehow. The backpropagation algorithm does just that, by exploiting the layered structure displayed in figure 1.

Before we can embark on deriving the algorithm we will introduce some notation. We assume L total layers, indexed as $l = 1, \dots, L$. Next we need to index the weights, nodes and biases. Let w_{jk}^l be the weight connecting k -th neuron in layer $l - 1$ and j -th neuron

in layer l . The index order in j and k are such that we can do matrix multiplication with index notation later down the road. Further let b_j^l be the bias for the j -th neuron in layer l . Thus the activation of the j -th neuron in layer l (a_j^l) becomes

$$a_j^l = \sigma \left(\sum_k w_{jk}^l a_k^{l-1} + b_j^l \right) = \sigma(z_j^l), \quad z_j^l \equiv \sum_k w_{jk}^l a_k^{l-1} + b_j^l. \quad (2.24)$$

Here σ is an activation function.

Now the cost function will depend directly on the activation of the output layer (a_j^L). However the activation of the output layer depends on the previous layers, meaning that the cost function depends indirectly on all the previous layers. Lets define the error Δ_j^L of the j -th neuron in layer L , as the change in cost function with respect to z_j^L .

$$\Delta_j^L \equiv \frac{\partial C}{\partial z_j^L} \quad (2.25)$$

We can similarly define the error of neuron j in layer l , as the change in the cost function with respect to z_j^l ,

$$\Delta_j^l \equiv \frac{\partial C}{\partial z_j^l} = \frac{\partial C}{\partial a_j^l} \frac{\partial a_j^l}{\partial z_j^l} = \frac{\partial C}{\partial a_j^l} \frac{d\sigma(z_j^l)}{dz_j^l}. \quad (2.26)$$

In the next few lines we are going to derive several equations needed for the algorithm, it will be apparent why after we have found them. Notice that (2.26) also can be written as

$$\Delta_j^l = \frac{\partial C}{\partial z_j^l} = \frac{\partial C}{\partial b_j^l} \frac{\partial b_j^l}{\partial z_j^l} = \frac{\partial C}{\partial b_j^l}. \quad (2.27)$$

Because $\partial b_j^l / \partial z_j^l = 1$ from (2.24). Again using the chain rule we can rewrite (2.26)

$$\begin{aligned} \Delta_j^l &= \frac{\partial C}{\partial z_j^l} = \sum_k \frac{\partial E}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} = \sum_k \Delta_k^{l+1} \frac{\partial z_k^{l+1}}{\partial z_j^l} \\ &= \left(\sum_k \Delta_k^{l+1} w_{kj}^{l+1} \right) \frac{d\sigma(z_j^l)}{dz_j^l}. \end{aligned} \quad (2.28)$$

To find the last equation, we differentiate the cost function with respect to the weight w_{jk}^l

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l} = \Delta_j^l a_k^{l-1}. \quad (2.29)$$

The equations (2.26), (2.27), (2.28) and (2.29) define what we call the backpropagation algorithm. Then, what exactly is the algorithm? It entails six steps:

- 1 Activation:** First activate the neurons in the activation layer (a_j^1) with desired data.
- 2 Feedforward:** Activate the nodes in following layers, this is done by equation (2.24).
- 3 Error at layer L :** Calculate the error at the last layer using (2.26).

4 Backpropagate error: With (2.28) we can calculate the error, iterating backwards in the network.

5 Calculate gradient: Find the gradient by using equations (2.27) and (2.29).

6 Update parameters: Update the parameters similarly to (2.5), however β_t are our weights and biases in this case.

The expression for updating the weights and biases are

$$w_{jk}^l \leftarrow w_{jk}^l - \eta \Delta_j^l a_k^{l-1} \quad (2.30)$$

$$b_j^l \leftarrow b_j^l - \eta \Delta_j^l \quad (2.31)$$

2.3.5 Initialization of weights

We mentioned earlier that the network has weights and biases between the layers. These needs to be initialized in some way. The biases are simple to initialize, as they are a single number for every node. These are initialized as a small, non-zero value b_0 , which we choose as $b_0 = 0.01$.

Before 2006, most neural networks were performing quite badly on most tasks, as they did not learn during training. One of the (several) reasons were due to bad initialization of weights. A common way of doing this was using the standard normal distribution $W_{i,j} \sim \mathcal{N}(0, 1)$. The problem with this is that it does not consider the size of the layers. In 2010, it was shown that when using Sigmoid as the activation function, Xavier-initialization give better results [1]. This is given as $W_{i,j} \sim \mathcal{U}\left(-\frac{1}{\sqrt{n}}, \frac{1}{\sqrt{n}}\right)$, where \mathcal{U} is the uniform distribution, and n is the number of nodes in the preceding layer.

In 2015, He-initialization was shown to work well with ReLU and Leaky ReLU [2]. Here, the weights are initialized using the normal distribution, but with a variance given by $v = 2/(1 + \alpha^2)n$, where again n is the number of nodes in the preceding layer, and α is the parameter of the Leaky ReLU-function. For ReLU, this is 0. This initialization only considers the number of nodes in the preceding layer, though normalized initialization considering also the the number of nodes in the succeeding layer could yield better results. We use the initialization given here, for the appropriate activation-function, and will not study the particular effects of this in depth.

3 Methods

As we mentioned in the introduction, we wish to study different ways of fitting two types of datasets. The first is fitting the *continuous* Franke Function (3.1), i.e. regression. We will use both stochastic gradient decent and feed forward neural network to fit the data. Both methods are covered in the theory sections 2.1 and 2.3 respectively.

Next we will embark on an classification problem. Namely predicting if breast tumors are malignant or benign, using the Breast Cancer Wisconsin data set². We will again use our

²<https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

feed forward neural network, as well as logistic regression. The latter is covered in the theory section 2.2.

3.1 Franke Function

$$f(x, y) = \frac{3}{4} \exp\left(-\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4}\right) + \frac{3}{4} \exp\left(-\frac{(9x+1)^2}{49} - \frac{(9y+1)^2}{10}\right) + \frac{1}{2} \exp\left(-\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4}\right) - \frac{1}{5} \exp(-(9x-4)^2 - (9y-7)^2). \quad (3.1)$$

In [4], we already studied the Franke function using linear regression, specifically OLS and OLS with an L2 and L1 parameter λ , so-called Ridge and Lasso regression. The results from these methods will form the basis for comparing our results using SGD and neural networks. The report, along with the code can be found at our GitHub³. In order to have comparable results, we will use the same parameters for the data. Only the methods will be different. In that project we generated the data using $N = 30 \times 30$ uniformly distributed datapoints in x - and y -direction, respectively. To simulate it being real data, we also added normally distributed noise with mean zero and standard deviation 0.2: $\epsilon \sim \mathcal{N}(0, 0.2)$. We also split the input and target data in the same way as before, using 80 % of the data for training and 20 % of the data for testing. The two splitted data sets are then scaled by subtracting the mean of the relevant training data.

3.1.1 Stochastic Gradient Descent

As in [4], we have to choose a model to fit the data to, when using SGD. The simplest is a bi-variate polynomial of degree P , such that our model will have p features. This is the design matrix X used in the previous project. Having obtained good results for OLS using $P = 6$, we use the same polynomial degree for SGD. Writing our own code for implementing SGD, we will analyze the results with MSE (equation (2.15)) as our cost function for various parameters. We include the L_2 regularization term, equation (2.17), in the gradient of the cost function, which gives the gradient corresponding to Ridge regression. The expression for the stochastic gradient is given in equation (3.2)

$$\nabla_{\beta} C^{\text{MB}}(\beta) = 2X^T(X\beta - \mathbf{z})/M + 2\lambda\beta \quad (3.2)$$

where OLS regression is obtained by setting $\lambda = 0$. There are multiple parameters to consider for stochastic gradient descent, and the ones we will study are the choice of learning rate, η , number of epochs, N_e , number of minibatches, N_B , regularization parameter, λ and the momentum parameter γ . This gives us a 5-dimensional hyperparameter space, and we would ideally optimize each parameter to get the lowest possible MSE. This could be done by gridsearch, where we test a range of values for every parameter. Testing 10 values for

³<https://github.com/sigurdru/FYS-STK4155/tree/main/project1>

each will mean that we have to run the algorithm 10^5 times, which is extremely slow. We will therefore simplify our search, only doing gridsearch over 1 or 2 of the parameters at a time, keeping the rest constant. We then study how each parameter affects the results, independently of the other parameters.

The optimal value for a given parameter might not be the optimal using different values of the other parameters. This method is then not guaranteed to optimize the entire parameter space. However, we can assume it will be a good approximation, and will save a great deal of time. Further, it is not too valuable to find the absolutely optimal parameters in our hyperparameter space, because there are many more parameters we will not controll for, like the noise in the data or polynomial degree of our model.

We begin by studying the MSE of the Franke function. For the MSE values we will include the result from both the train and test data in our first simulation only, in order to confirm that overfitting takes place eventually, while focusing on the test MSE in all the remining simulation, which is the actual quantity of interest. The MSE is studied for different values of η as a function of $N_e \in [0, 150]$, setting $\lambda = \gamma = 0$ and $N_B = 20$, i.e. minibatches of size $M = 36$ for the 720 points in the training dataset. We choose 21 evenly spaced values of $\eta \in [0.01, 0.9]$.

We then choose one of the favorable learning rates to study the MSE for minibatch sizes of $M \in [720, 360, 240, 144, 72, 48, 36, 30, 24]$, as a function of epochs, using $N_e \in [0, 150]$ once again. The values of M are chosen such that N_B becomes integers. This does not have to be the case, but will make the results more consistent. Next, using $N_B = 20$ minibatches, we then study the MSE after $N_e = 150$ epochs as a function of η and λ to study the effect of regularization. For this we use 11 linearly distributed values of $\eta \in [0.1, 0.5]$ and 11 logarithmically distributed values of $\lambda \in [10^{-5}, 1]$. We only consider small λ values, since our previous study of the Franke function with linear regression indicated that higher λ yielded poor results [4]. Finally, we study the MSE over 150 epochs for 15 evenly spaced values of $\gamma \in [0, 0.7]$ with a fixed learning rate.¹

We expect the results to be very dependent on the choice of lerning rate η . It is important to not choose too low, such that we actually train, and not too high, such that we can not reach a minimum. A way of balancing this is using a dynamic learning rate, which starts high, such that the model gets better quick in the beginning, and then decreases as function of epoch, preventing the model from overshooting and circle around the minimum in β -space. We choose the following function for dynamic learning rate, where $t \in [0, N_e]$ denotes the current epoch.

$$\eta_t = \eta_0 \cdot \left(1 - \frac{t}{N_e}\right) \quad (3.3)$$

It will decrease lineraly from η_0 , to $\eta_{N_e} = 0$ at the last epoch. We could have chosen other shapes for the dynamic learning rate, for instance exponential decreasing η , or an upside-down sigmoid. It can be useful to study the effects of different learning schedules, though this is outside the scope of this paper.² Using (3.3), we remake the first figure, to

¹ Burde
man kan-
skje lagd
tabell
med
paramete-
terne?

² Dette
burde
kanskje
stå i
duskisjon

study the effects of changing η during training.

After the SGD has been studied and appropriate parameters have been estimated, we will plot the resulting fit of the gradient descent using the optimal parameters found, and compare it to the Franke function as an indication of the result. This is a sanity check to ensure that the model actually makes sense. Although the fit was performed on a data set with noise $\epsilon \sim \mathcal{N}(0, 0.2)$ we will use $\epsilon \sim \mathcal{N}(0, 0.05)$ when plotting the prediction in order to see the details of the surface more clearly. ³

³ *MÅ
GJØRES!*

3.1.2 Feed Forward Neural Network

When using the neural network to fit the Franke function, we use a lot of the same methods as for SGD. A key difference is that instead of iteratively updating a model β , we now train a network of several layers, each with many nodes. One of the results of this is that we do not have to choose the shape the model will take.

Instead of giving our network the design matrix X for a certain polynomial degree P , we can pass it only the collection of points (x_i, y_i) , and let the network adjust the weights and biases accordingly. Since the Franke function is an exponential function we know that it can be approximated as a higher order polynomial, so by using the design matrix as an input we exploit this property such that the network converges faster. Having already fitted the Franke function with a design matrix with linear regression and SGD, we now choose the x_i and y_i values only. Not providing the network with any initial information has the advantage that we get a more rigorous test of the network's performance, since we ensure that the result is not directly reliant on the information in question. Another important motivation for this is that if we were to fit some other data, e.g. terrain data, we may not have any a-priori information regarding the input data. Omitting the design matrix when we train our neural network will thus yield a final model capable of fitting various types of data.

Since we are dealing with a regression problem and we're fitting a continuous function, a natural choice of the cost function is the MSE. For the neural network we will not compute the total MSE of the output layer as we have previously done, but the individual MSE of each output node. This takes into account the error at each individual output neuron when we update the weights with backpropagation.

As with SGD, neural networks have many parameters we have to tune. To reduce the dimensionality of the hyperparameter space, we choose some to study, namely the learning rate η , number of epochs N_e , regularization parameter λ , momentum parameter γ , and choice of activation function. We will also study the number of hidden layers and number of nodes in these. We then do not change the batch size, choosing one based on the results from SGD, nor the weight initialization, using the ones described in section 2.3.5, given the activation function. This is still a very big space, so again we are more interested in how the different parameters affect the result independently of the others.

A good indicator for our neural network being implemented correctly is making sure that it learns during training. This we check by plotting the MSE calculated after every epoch.

We do this for different η s. We will then also find reasonable choices for the learning rate. We choose two hidden layers with ten nodes in each, and use the sigmoid activation function. As mentioned, as we are fitting a function, there is no activation function for the output layer. For the initial analysis we again choose $\lambda = \gamma = 0$.

Having initialized the neural network we are now going to train it. Each training iteration begins by randomly shuffling the data and dividing them into minibatches, just as we did when we performed the SGD analysis. Now the feedforward and backpropagation are performed for each minibatch.

For the first simulation we choose 21 linearly distributed values of $\eta \in [10^{-3}, 0.9]$, and plot the result MSE from the training data and testing data. We will then repeat this analysis using 21 new η values on an interval where the resulting MSE was low, plotting the test MSE only, like in the previous section. We also do this for the R^2 score, given in equation (3.4), where \mathbf{z} and $\tilde{\mathbf{z}}$ is the target and prediction values, respectively. This is because the R^2 -score gives a relative performance, while MSE as an absolute measure. More discussion on these can be found in [4].

$$R^2(\mathbf{z}, \tilde{\mathbf{z}}) = 1 - \frac{\sum_{i=0}^{n-1} (z_i - \tilde{z}_i)^2}{\sum_{i=0}^{n-1} (z_i - \bar{z})^2}, \quad \bar{z} = \frac{1}{n} \sum_{i=0}^{n-1} z_i \quad (3.4)$$

We then look at the MSE using dynamic learning rate, keeping the other parameters the same. Next the activation function for the hidden layers is tested using RELU and leaky RELU instead of sigmoid. We also study the dependence on the regularization parameter λ . Choosing ⁴

Lastly, we study the dependence of different numbers of layers and nodes, using 1 layer with 10 nodes, 2 with 20 each and 3 with 20 nodes in first

⁴ må lage
plot for
dette

So far, we have only studied the statistical accuracy of our neural network. As a sanity check we will use the optimal parameters obtained in order to plot the resulting prediction of the neural network, using $\epsilon \sim \mathcal{N}(0, 0.05)$ just as we did for SGD. Some of the parameters we use may not be ideal when the noise is reduced, but the purpose of this test is simply to see whether our model behaves in a desired way, so reducing the noise makes the comparison easier.

3.2 Wisconsin Breast Cancer Data

3.2.1 Feed Forward Neural Network

3.2.2 Logistic Regression

Lastly we want to study the breast cancer data using logistic regression. See theory section 2.2 for an explanation of the method. We use the Sigmoid (2.11) as our soft classifier in the output, and cross entropy (2.8) (gradient given by (2.9)) as our cost function. See theory section 2.3.3 for derivation of the cross entropy. To optimize the weight and biases

we implement stochastic gradient decent algorithm (2.5), using five minibatches. We will also study the implementation of L_2 -regularization (2.17). Our aim is to compare the result with those obtained from out feed forward neural network and Scikit-Learn’s logistic regression functionality.

The main quantity we will study is the accuracy score, which is the number of correct classification divided by total cases. The Sigmoid is not a hard classifier, i.e. it returns continuous values between zero and one. To calculate the accuracy we will therefore round the answer such that any output less than 0.5 will be set to zero, and above set to one. This means that we loose any interpretation of how *sure* our algorithm is of the result. For example if our output is 0.51, then it will set the result to 1, even though that could be interpreted as inconclusive. Thus we will perform a small analysis in the end on how confident our algorithm is. More precisely we will, for what we interpret as the optimal parameters, look at how many cases are classified within the range (0.1,0.9). We discuss this a bit more in-depth in the discussion-section.

We initialize the weights and bias with random normal distribution with mean around zero and standard deviation of one. We want to look at what happens to the accuracy as epochs increases, for different learning rates. Therefore we first plot the accuracy score as a function of learning rate and number of epochs. Specifically we will use 200 epochs, and 10 different learning rates η between 10^{-5} and $10^{-0.5}$ distributed evenly on a logarithm scale. We expect to see overfitting when η is large, and underfitting when η is small.

Next we are ready to perform an grid search, to find the optimal learning rate η and L_2 parameter λ , still using 200 epochs. Therefore we again plot the accuracy score, however this time as a function of λ and η , picking out the accuracy score after the last epoch. Specific parameters we will test for are 10 different learning rates η and regularization parameters λ , distributed evenly on a logarithm scale. Learning rate will be between 10^{-5} and 10^{-1} , and λ will be between 10^{-5} and 10^{-3} .

4 Results

4.1 Franke Function

4.1.1 Stochastic Gradient Descent

We plot the MSE of the Franke function using SGD for our training data and test data for $\eta \in [0.01, 0.9]$, shown on the left and right panel of figure 3 respectively. The MSE was initially much higher than the maximum colorbar value of 0.07, but we have chosen this as an upper limit, as it provided much more insight regarding nuances of the MSE. From our previous study of the Franke function with OLS regression [4], we found that the MSE was below 0.07 for polynomial degrees $P \in [1, 8]$ for the test data with 90 bootstrap iterations. A maximum value of MSE=0.07 is thus a natural choice for the desired result.

In figure 3 we see that the test MSE converges faster for increased values of η as expected. For $\eta = 0.01$ we get an MSE barely below 0.07 after $N_e = 150$, as the gradient steps are too





Figure 3. Initial MSE of the train and test data as a function of epochs for different learning rates η . A maximum MSE value of 0.07 is set, so dark red regions may correspond to significantly higher MSE than it appears to.

small. When η exceeds 0.5 we get abrupt increments of the test MSE after certain epochs. These learning rates are too high, as a gradient step goes beyond the actual minima. For $\eta = 0.9$ this effect is apparent, with large fluctuations of the MSE over certain epochs. We emphasize that the bright red lines present could represent MSE values far above 0.07.

We proceed by plotting the test MSE for $\eta \in [0.01, 0.5]$, shown in figure 4, where we set the maximum MSE to 0.07 once again. We clearly see how the increased learning rate yields a faster converging result, but at the expense of the MSE stability, which we can see from clear fluctuations for $\eta = 0.5$. When our prediction approaches the desired model, the gradients are relatively small, such that high learning rates potentially overshoots minima.

For studying the MSE for different number of minibatches and different momentum parameters, γ , we will choose $\eta = 0.25$. In figure 4, we see that $\eta = 0.25$ yields fairly quick convergence of the MSE without significant variations. Learning rates higher than this appears to be too high, while the ones below it appears to be small enough that minima of the cost function are rarely overshoot. $\eta = 0.25$ appears therefore to be a boundary between slow convergence and minima overshooting, and when testing other parameters we are more likely to encapsule both cases at once. *Kanskje skrive at flere etaer gir gode resultat, og at vi velger 0.25 fordi vi er mer interessert i å se på hvordan andre parametere påvirker resultat uavhengig, og at 0.25 derfor er et ok valg? Jeg (HO) synes ikke det er serlig klart at høyere eta er for ustabile. Er kanskje en endre grense, men absolut ikke en øvre.*




Figure 4. Test MSE for different η values. $\eta < 0.25$ gives a relatively slowly converging MSE, while $\eta > 0.25$ gives fluctuating MSE values. Maximum MSE= 0.07 is chosen.

The MSE as a function of epochs for different number of minibatches is shown in figure 5. We use $\text{MSE}_{\max} = 0.07$ once again. Figure 5 shows an expected behaviour of the MSE.

By not using enough minibatches we are unable to get an MSE below 0.06 after $N_e = 150$, which is most likely due to model being stuck in a local minimum substantially worse than the global one, unable to get out. When we use $N_B = 30$ we get a quickly converging MSE with a low minimum value, but there appears to some fluctuations. Because of the small batch size, the batches are more prone to consist of points not giving a good approximation of the gradient of the whole dataset. ⁵




Figure 5. Test MSE for different number of minibatches. The bottom line corresponds to no minibatches overall, and clearly gives the worst result. The heatmap is produced with $\text{MSE}_{\max} = 0.07$.

We plot the test MSE after $N_e = 150$ epochs with different regularization parameters, shown in figure 6. Using the same η interval as before and $\lambda \leq 1$ we do not set an upper limit on the shown MSE, as we are interested in the general impact of the λ values, and not seeking nuances of the lowest values as before. We see in the figure that the results are not very affected by the choice of η . This is because we only see the last MSE of the training. However, the MSE is lowest for small values of λ , increasing quickly for higher values. This is similar to what we saw in [4], though then we observed a small dip in the MSE for $\log_{10}(\lambda) = -2.07$ for certain polynomial degrees.

Figure 7 shows the resulting MSE values for the test data with different momentum parameters as a function of $N_e \in [0, 150]$. We see that increasing γ values causes the MSE to drop faster, particularly in the beginning when the gradients are steep, which is expected. We also notice that $\gamma \gtrsim 0.35$ results in fluctuations of the MSE, especially evident for $\gamma = 0.7$, where it appears that step sizes taken following large gradient values are too big. It appears that $\gamma = 0.2$ is the optimal value for this particular test, as it is small enough to avoid overshooting minima for high gradients, while still increasing the convergence rate. However, the overall impact is not very significant.




Figure 6. MSE for different values λ , corresponding to Ridge regression, for different η values.




Figure 7. MSE evolution for different momentum parameters, using $\text{MSE}_{\max} = 0.07$, as before. The bottom line corresponds to no momentum.

The MSE calculated with dynamic η is shown in figure 8, using $\eta_0 \in [0.1, 0.9]$. We include the η values previously found to be too high, in order to see whether the scheduled learning gives better results here, while omitting values of $\eta < 0.1$. The figure gives us the expected

result that high initial values of η no longer causes abrupt variations of the MSE towards the end of the simulation, as $\eta \rightarrow 0$ and $N_e \rightarrow 150$. Equation (3.3) appears to be a very good choice for scaling the learning rate when doing SGD, since all values of $\eta_0 > 0.5$ gives decent results after 150 epochs. These solutions converge quickly in the beginning, and as the learning rate drops we get very stable evolutions of the MSE, with no apparent instability in MSE. It appears however that the learning rates become so small towards the end that there is little evolution of the MSE after 120 epochs, meaning that we might not actually have reached an optimal MSE score if the small learning rates near the end have effectively halted convergence. Further, the dynamic learning schedule gives more or less identical MSE values in the end for $\eta_0 > 0.5$. The advantage of this is that our model is less dependant on very specific parameter choices for η . If however, the final result is inadequate, adjusting the initial learning rate will no longer have a desired impact. More training epochs is then required to get optimal results.

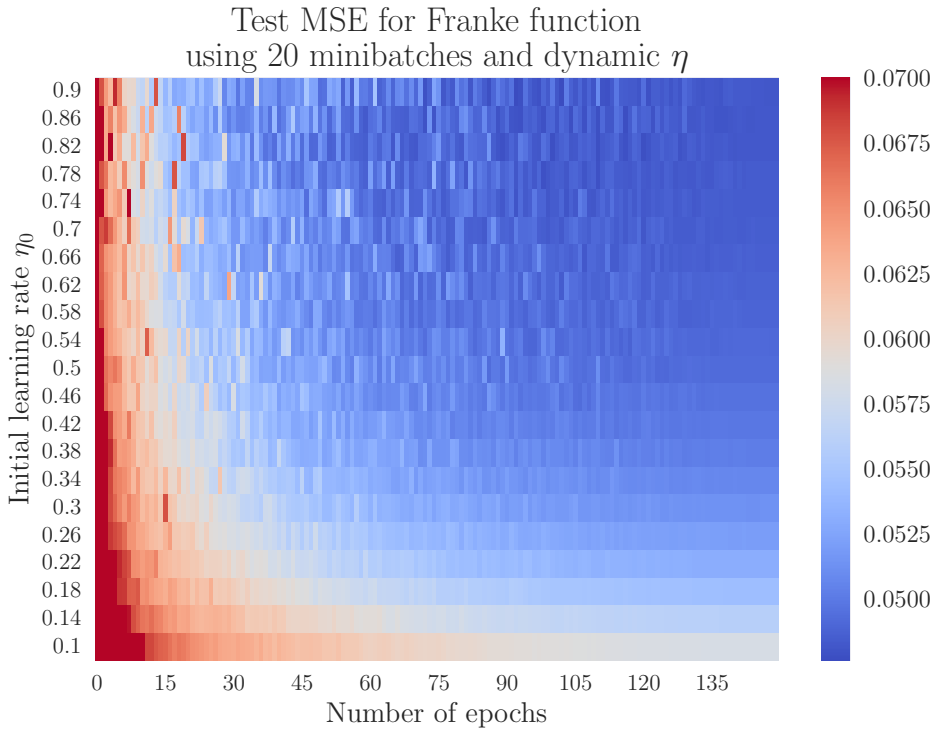


Figure 8. Evolution of the MSE for different values of η_0 , all of them decreasing linearly to a final value of $\eta = 0$ after 150 epochs.

4.1.2 Feed Forward Neural Network

We plot the train and test MSE of the Franke function for different learning rates as a function of epochs, shown in the upper and lower panel of figure 9 respectively. We do not cap the upper limit of the MSE. Even though the two plots look similar, the value of the MSE is higher for the test data, which is as expected. We will in the following figures cap the MSE at 0.07, like what was done for SGD.

We choose $\eta \in [0.05, 0.5]$ and plot the test MSE as a function of $N_e = 250$ epochs shown in figure 10. The figure shows how the MSE drops as we continue to train the network, with little evidence of overfitting, which would be a general increases of the MSE towards the





Figure 9. MSE of the train and test data after 150 epochs for different learning rates using a Neural network.

end of training. We do however see some clear vertical lines; abrupt increases of the MSE for all η . These were not expected, though can possibly be explained through the way the

network is trained. In order to keep results comparable over many runs, the seed of the random distribution is set before each net is created. This was mainly to ensure that the weights are initialized the same for all. However, a side effect of this is that the dataset is split into the same batches for each net, at a given epoch. Sometimes these batches lead to a bad approximation of gradient for the entire dataset, giving worse test results. Here that happens at the same epoch every time. ⁶

6 skrive
noe om
hvorfor
vi ikke
endrer
det?




Figure 10. Test MSE of the Franke function for different learning rates.

We proceed by plotting the R^2 score, setting a minimum value of $R^2 = 0.55$ to highlight results in the same vicinity of what we got for the linear regression study, shown in figure 11. The figure shows that we obtain R^2 scores exceeding the ones we obtained with linear regression, where the test $R^2_{\text{test}} < 0.6$ was obtained without resampling, while we exceed this value for multiple configurations of the neural network. Other than numerical results, the R^2 score does not appear to provide any new information, as its development essentially displays inverse properties of the MSE, where high values correspond to low MSE values and vice versa. The MSE increments at certain epochs is now present in the form of drops in the R^2 scores.

Figure 12 shows the evolution of the MSE when we use dynamic learning rates, calculated with equation (3.3) as before. The result is similar to what we obtained with SGD, with suppressed fluctuations of the MSE for higher epochs. This has a fairly large effect on the Neural network, as the fluctuations of the MSE for $N_e > 200$ in figure 10 are barely visible. Once again, our scaling algorithm have provided MSE values towards the end of the simulation that varies little for different learning rates.




Figure 11. The R^2 score of the Neural network, using a minimum value of the heatmap of $R^2_{\min} = 0.55$.

As a final test of our neural network, we plot the test MSE over 250 epochs using the Relu and Leaky Relu activation functions for the hidden layers, shown in the upper and lower panel of figure 13 respectively. An upper value of 0.07 have been used for the MSE once again. Values of $\eta > 0.15$ yielded poor results for these activation functions, and caused the network to not converge for some of the higher learning rates. The behaviour of the MSE for these two activation functions differ very little, but it appears that Leaky Relu yields somewhat lower MSE values than Relu. We include the corresponding MSE evolution for the same η values with sigmoid function as activation for the hidden layers, shown in figure 14.

Figure 14 emphasizes the vast difference of the neural network with different activation functions for the hidden layers. Both Relu and Leaky Relu converges a lot faster than sigmoid activation, and while there seems to be slight increase of MSE values towards the end for the former two, these are less present for the lower values of η , which might have been a more suitable choice. Other than that, there are still fluctuations for all three activation functions, and the sigmoid function converges much slower than the other two. Figure 14 shows that the MSE irregularities in figure 10 are still present, although barely visible the the learning rates where the MSE barely drops below 0.06. The MSE irregularities near the end are much smaller for the other two activation functions, and doesn't seem to affect the lowest learning rates to a noticeable degree.

We now plot different network configurations. HÅKON OLAV




Figure 12. MSE of the Franke function from the Neural network, using a linearly decreasing learning rate, where the initial learning rates are shown on the y -axis.

4.2 Wisconsin Breast Cancer Data

4.2.1 Feed Forward Neural Network





Figure 13. MSE Relu and Leaky Relu




Figure 14. MSE of the Neural network with the sigmoid activation function for the hidden layers. The learning rates and epochs are identical to that of figure 13, allowing us to compare the effect of the different activation functions.

4.2.2 Logistic Regression

In figure 15 we plot the accuracy score as a function of epochs, for different learning rates η and L_2 parameter set to $\lambda = 0$. The top plot is for the train data, and bottom for test data. As mentioned we want to see how the network learns, meaning that we are not that interested in the precise accuracy score before we perform a grid search. First, looking at the train accuracy, when $\eta \leq 3 \cdot 10^{-5}$ we do not converge within the 200 epochs (underfitting). We also notice faint signs of overfitting for $\eta \geq 0.1122$. These effects is more apparent in the test accuracy, where we see clear overfitting at $\eta \geq 0.03981$, and underfitting at $\eta \leq 8 \cdot 10^{-5}$. The optimal learning rate should therefore be within $8 \cdot 10^{-5} \geq \eta \leq 0.1122$.

In figure 16 we have plotted the results from our grid search. There we have displayed the precise test accuracy score, as a function of learning rate η and hyper parameter λ . Looking at the figure, we notice that our best accuracy score is 99.1%, for $0.001 \geq \eta \leq 0.00631$ and all λ 's. We acquire the same accuracy score for $\eta = 0.01585$ and $\lambda \geq 10^{-3.6}$. This is the exact same accuracy as when using Scikit-Learn's logistic regression with L_2 regularization, and corresponds to one miss-categorization.

There are many potential optimal parameters, however lets choose $\eta = 0.00251$ and $\lambda = 0$, and print the number of cases where our output is withing (0.1, 0.9). We remind that this is not a very comprehensive analysis, and we talk more about this in the discussion-section. Nevertheless, after the first epoch we have 58 cases, and after 200 epochs we have eight cases where the output is within (0.1, 0.9).


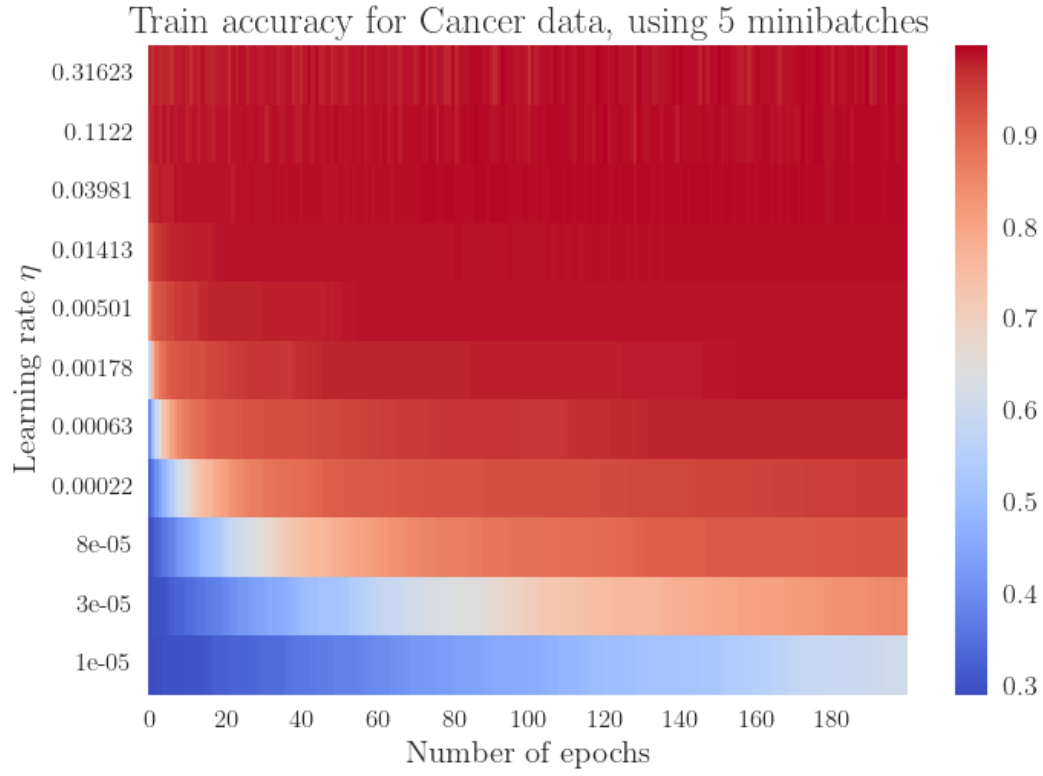


Figure 15. These two figures illustrates how the accuracy score evolves as a function of epochs (x -axis) and learning rate η (y -axis). An accuracy score of one corresponds to 100% correct classifications.

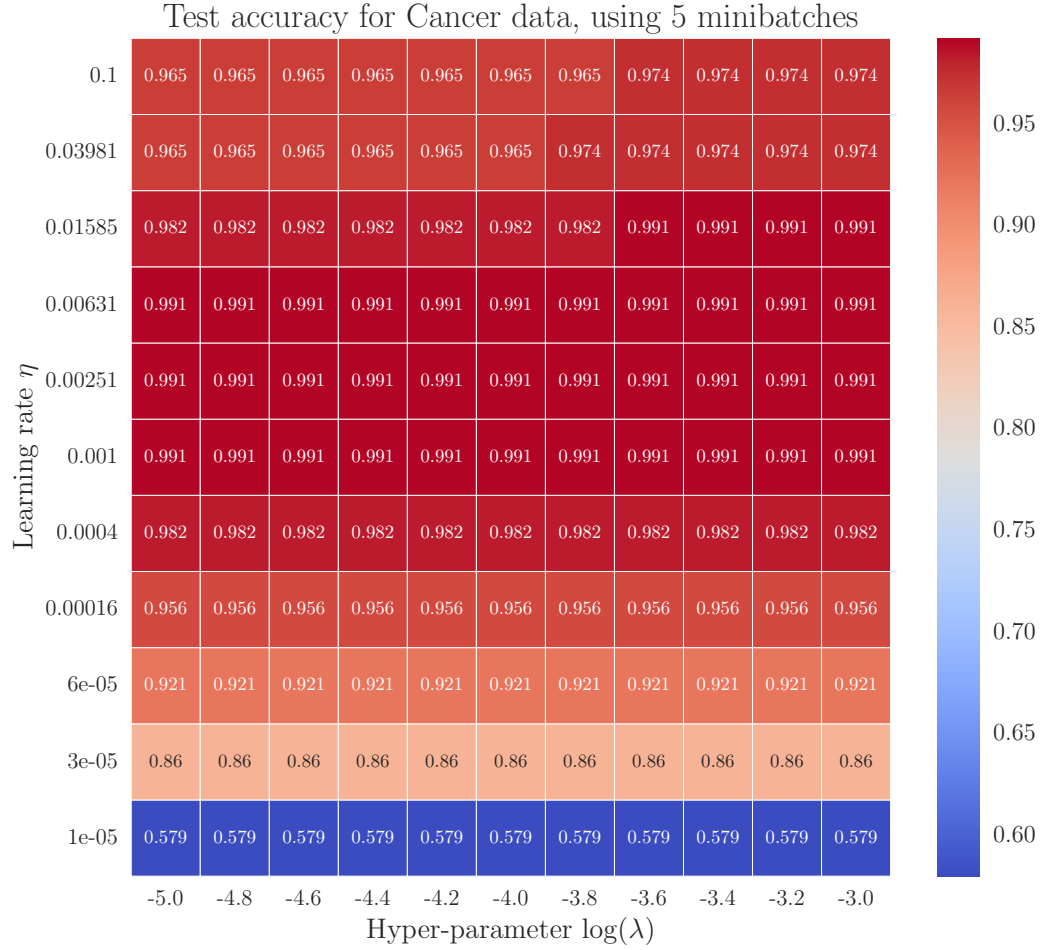


Figure 16. Here we have plotted the accuracy for the test data, as a function of both learning rate η (y -axis) and logarithm of hyper parameter λ (x -axis). An accuracy score of one corresponds to 100% correct classification.

5 Discussion

LOGISTIC REGRESSION:

The fact that we obtain the same accuracy (99.2%) as when using Scikit-Learn’s logistic regression, gives us confidence in our results. The 99.2% accuracy corresponds to one misclassification. We do not believe it feasible to better this result, given that a professionally written library yields the same. One area which we would have liked to analyze further is the confidence of our algorithm. After 200 epochs, for $\eta = 0.00251$ and $\lambda = 0$, we had eight cases where the output was within (0.1, 0.9). That threshold is arbitrary, however it illustrates potential further analysis. To find the optimal parameters we should not only analyze accuracy, but also minimize the number of cases which one deems to inconclusive. Further analysis one could try, is implementing different output functions, initialization of weights and bias and momentum parameter (2.6).

6 Conclusion

References

- [1] Xavier Glorot and Y. Bengio. Understanding the difficulty of training deep feedforward neural networks. *Journal of Machine Learning Research - Proceedings Track*, 9:249–256, 01 2010.
- [2] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. *IEEE International Conference on Computer Vision (ICCV 2015)*, 1502, 02 2015.
- [3] Pankaj Mehta, Marin Bukov, Ching-Hao Wang, Alexandre G.R. Day, Clint Richardson, Charles K. Fisher, and David J. Schwab. A high-bias, low-variance introduction to machine learning for physicists. *Physics Reports*, 810:1–124, May 2019.
- [4] Håkon Olav Torvik, Vetle Vikenes, and Sigurd Sørle Rustad. Analysis of regression and resampling methods. *University of Oslo*, October 2021.